
[Prev](#)[Next](#)

Chapter 3. Configuration

[3.1. Programmatic configuration](#)

[3.2. Obtaining a SessionFactory](#)

[3.3. JDBC connections](#)

[3.4. Optional configuration properties](#)

[3.4.1. SQL Dialects](#)

[3.4.2. Outer Join Fetching](#)

[3.4.3. Binary Streams](#)

[3.4.4. Second-level and query cache](#)

[3.4.5. Query Language Substitution](#)

[3.4.6. Hibernate statistics](#)

[3.5. Logging](#)

[3.6. Implementing a NamingStrategy](#)

[3.7. XML configuration file](#)

[3.8. J2EE Application Server integration](#)

[3.8.1. Transaction strategy configuration](#)

[3.8.2. JNDI-bound SessionFactory](#)

[3.8.3. Current Session context management with JTA](#)

[3.8.4. JMX deployment](#)

Hibernate is designed to operate in many different environments and, as such, there is a broad range of configuration parameters. Fortunately, most have sensible default values and Hibernate is distributed with an example `hibernate.properties` file in `etc/` that displays the various options. Simply put the example file in your classpath and customize it to suit your needs.

3.1. Programmatic configuration

An instance of `org.hibernate.cfg.Configuration` represents an entire set of mappings of an application's Java types to an SQL database. The `org.hibernate.cfg.Configuration` is used to build an immutable `org.hibernate.SessionFactory`. The mappings are compiled from various XML mapping files.

You can obtain a `org.hibernate.cfg.Configuration` instance by instantiating it directly and specifying XML mapping documents. If the mapping files are in the classpath, use `addResource()`. For example:

```
Configuration cfg = new Configuration()
    .addResource("Item.hbm.xml")
    .addResource("Bid.hbm.xml");
```

An alternative way is to specify the mapped class and allow Hibernate to find the mapping document for you:

```
Configuration cfg = new Configuration()
    .addClass(org.hibernate.auction.Item.class)
    .addClass(org.hibernate.auction.Bid.class);
```

Hibernate will then search for mapping files named **/org/hibernate/auction/Item.hbm.xml** and **/org/hibernate/auction/Bid.hbm.xml** in the classpath. This approach eliminates any hardcoded filenames.

A `org.hibernate.cfg.Configuration` also allows you to specify configuration properties. For example:

```
Configuration cfg = new Configuration()
    .addClass(org.hibernate.auction.Item.class)
    .addClass(org.hibernate.auction.Bid.class)
    .setProperty("hibernate.dialect", "org.hibernate.dialect.MySQLInnoDBDialect")
    .setProperty("hibernate.connection.datasource", "java:comp/env/jdbc/test")
    .setProperty("hibernate.order_updates", "true");
```

This is not the only way to pass configuration properties to Hibernate. Some alternative options include:

1. Pass an instance of `java.util.Properties` to `Configuration.setProperties()`.
2. Place a file named **hibernate.properties** in a root directory of the classpath.
3. Set System properties using `java -Dproperty=value`.
4. Include `<property>` elements in `hibernate.cfg.xml` (this is discussed later).

If you want to get started quickly **hibernate.properties** is the easiest approach.

The `org.hibernate.cfg.Configuration` is intended as a startup-time object that will be discarded once a `SessionFactory` is created.

3.2. Obtaining a SessionFactory

When all mappings have been parsed by the `org.hibernate.cfg.Configuration`, the application must obtain a factory for `org.hibernate.Session` instances. This factory is intended to be shared by all application threads:

```
SessionFactory sessions = cfg.buildSessionFactory();
```

Hibernate does allow your application to instantiate more than one `org.hibernate.SessionFactory`. This is useful if you are using more than one database.

3.3. JDBC connections

It is advisable to have the `org.hibernate.SessionFactory` create and pool JDBC connections for you. If you take this approach, opening a `org.hibernate.Session` is as simple as:

```
Session session = sessions.openSession(); // open a new Session
```

Once you start a task that requires access to the database, a JDBC connection will be obtained from the pool.

Before you can do this, you first need to pass some JDBC connection properties to Hibernate. All Hibernate property names and semantics are defined on the class `org.hibernate.cfg.Environment`. The most important settings for JDBC connection configuration are outlined below.

Hibernate will obtain and pool connections using `java.sql.DriverManager` if you set the following properties:

Table 3.1. Hibernate JDBC Properties

Property name	Purpose
<code>hibernate.connection.driver_class</code>	<i>JDBC driver class</i>
<code>hibernate.connection.url</code>	<i>JDBC URL</i>
<code>hibernate.connection.username</code>	<i>database user</i>
<code>hibernate.connection.password</code>	<i>database user password</i>
<code>hibernate.connection.pool_size</code>	<i>maximum number of pooled connections</i>

Hibernate's own connection pooling algorithm is, however, quite rudimentary. It is intended to help you get started and is *not intended for use in a production system*, or even for performance testing. You should use a third party pool for best performance and stability. Just replace the [hibernate.connection.pool_size](#) property with connection pool specific settings. This will turn off Hibernate's internal pool. For example, you might like to use c3p0.

C3P0 is an open source JDBC connection pool distributed along with Hibernate in the **lib** directory. Hibernate will use its `org.hibernate.connection.C3P0ConnectionProvider` for connection pooling if you set [hibernate.c3p0.*](#) properties. If you would like to use Proxool, refer to the packaged **hibernate.properties** and the Hibernate web site for more information.

The following is an example **hibernate.properties** file for c3p0:

```
hibernate.connection.driver_class = org.postgresql.Driver
hibernate.connection.url = jdbc:postgresql://localhost/mydatabase
hibernate.connection.username = myuser
hibernate.connection.password = secret
hibernate.c3p0.min_size=5
hibernate.c3p0.max_size=20
hibernate.c3p0.timeout=1800
hibernate.c3p0.max_statements=50
hibernate.dialect = org.hibernate.dialect.PostgreSQLDialect
```

For use inside an application server, you should almost always configure Hibernate to obtain connections from an application server `javax.sql.DataSource` registered in JNDI. You will need to set at least one of the following properties:

Table 3.2. Hibernate Datasource Properties

Property name	Purpose
hibernate.connection.datasource	<i>datasource JNDI name</i>
hibernate.jndi.url	<i>URL of the JNDI provider (optional)</i>
hibernate.jndi.class	<i>class of the JNDI InitialContextFactory (optional)</i>
hibernate.connection.username	<i>database user (optional)</i>
hibernate.connection.password	<i>database user password (optional)</i>

Here is an example **hibernate.properties** file for an application server provided JNDI datasource:

```
hibernate.connection.datasource = java:/comp/env/jdbc/test
hibernate.transaction.factory_class = \
    org.hibernate.transaction.JTATransactionFactory
hibernate.transaction.manager_lookup_class = \
    org.hibernate.transaction.JBossTransactionManagerLookup
hibernate.dialect = org.hibernate.dialect.PostgreSQLDialect
```

JDBC connections obtained from a JNDI datasource will automatically participate in the container-managed transactions of the application server.

Arbitrary connection properties can be given by prepending "hibernate.connection" to the connection property name. For example, you can specify a [charSet](#) connection property using [hibernate.connection.charSet](#).

You can define your own plugin strategy for obtaining JDBC connections by implementing the interface `org.hibernate.connection.ConnectionProvider`, and specifying your custom implementation via the [hibernate.connection.provider_class](#) property.

3.4. Optional configuration properties

There are a number of other properties that control the behavior of Hibernate at runtime. All are optional and have reasonable default values.

Warning

*Some of these properties are "system-level" only. System-level properties can be set only via `java -Dproperty=value` or **hibernate.properties**. They *cannot* be set by the other techniques described above.*

Table 3.3. Hibernate Configuration Properties

Property name	Purpose
hibernate.dialect	<p>The classname of a Hibernate <code>org.hibernate.dialect.Dialect</code> which allows Hibernate to generate SQL optimized for a particular relational database.</p> <p>e.g. <code>full.classname.of.Dialect</code></p> <p>In most cases Hibernate will actually be able to choose the correct <code>org.hibernate.dialect.Dialect</code> implementation based on the JDBC metadata returned by the JDBC driver.</p>
hibernate.show_sql	<p>Write all SQL statements to console. This is an alternative to setting the log category <code>org.hibernate.SQL</code> to debug.</p> <p>e.g. <code>true</code> <code>false</code></p>
hibernate.format_sql	<p>Pretty print the SQL in the log and console. e.g. <code>true</code> <code>false</code></p>
hibernate.default_schema	<p>Qualify unqualified table names with the given schema/tablespace in generated SQL. e.g. <code>SCHEMA_NAME</code></p>
hibernate.default_catalog	<p>Qualifies unqualified table names with the given catalog in generated SQL. e.g. <code>CATALOG_NAME</code></p>
hibernate.session_factory_name	<p>The <code>org.hibernate.SessionFactory</code> will be automatically bound to this name in JNDI after it has been created.</p> <p>e.g. <code>jndi/composite/name</code></p>
hibernate.max_fetch_depth	<p>Sets a maximum "depth" for the outer join fetch tree for single-ended associations (one-to-one, many-to-one). A 0 disables default outer join fetching.</p> <p>e.g. recommended values between 0 and 3</p>
hibernate.default_batch_fetch_size	<p>Sets a default size for Hibernate batch fetching of associations. e.g. recommended values 4, 8, 16</p>
hibernate.default_entity_mode	<p>Sets a default mode for entity representation for all sessions opened from this <code>SessionFactory</code></p> <p><code>dynamic-map</code>, <code>dom4j</code>, <code>pojo</code></p>
hibernate.order_updates	<p>Forces Hibernate to order SQL updates by the primary key value of the items being updated. This will result in fewer transaction deadlocks in highly concurrent systems. e.g. <code>true</code> <code>false</code></p>
hibernate.generate_statistics	<p>If enabled, Hibernate will collect statistics useful for performance tuning. e.g. <code>true</code> <code>false</code></p>
hibernate.use_identifier_rollback	<p>If enabled, generated identifier properties will be reset to default values when objects are deleted. e.g. <code>true</code> <code>false</code></p>

Property name	Purpose
hibernate.use_sql_comments	<p>If turned on, Hibernate will generate comments inside the SQL, for easier debugging, defaults to false.</p> <p>e.g. true false</p>

Table 3.4. Hibernate JDBC and Connection Properties

Property name	Purpose
hibernate.jdbc.fetch_size	A non-zero value determines the JDBC fetch size (calls <code>Statement.setFetchSize()</code>).
hibernate.jdbc.batch_size	<p>A non-zero value enables use of JDBC2 batch updates by Hibernate. e.g. recommended values between 5 and 30</p> <p>Set this property to true if your JDBC driver returns correct row counts from <code>executeBatch()</code>. It is usually safe to turn this option on. Hibernate will then use batched DML for automatically versioned data. Defaults to false.</p> <p>e.g. true false</p>
hibernate.jdbc.batch_versioned_data	<p>Select a custom <code>org.hibernate.jdbc.Batcher</code>. Most applications will not need this configuration property.</p> <p>e.g. <code>classname.of.BatcherFactory</code></p>
hibernate.jdbc.use_scrollable_resultset	<p>Enables use of JDBC2 scrollable resultsets by Hibernate. This property is only necessary when using user-supplied JDBC connections. Hibernate uses connection metadata otherwise.</p> <p>e.g. true false</p>
hibernate.jdbc.use_streams_for_binary	<p>Use streams when writing/reading binary or serializable types to/from JDBC. <i>*system-level property*</i></p> <p>e.g. true false</p>
hibernate.jdbc.use_get_generated_keys	<p>Enables use of JDBC3 <code>PreparedStatement.getGeneratedKeys()</code> to retrieve natively generated keys after insert. Requires JDBC3+ driver and JRE1.4+, set to false if your driver has problems with the Hibernate identifier generators. By default, it tries to determine the driver capabilities using connection metadata.</p> <p>e.g. true false</p>
hibernate.connection.provider_class	<p>The <code>classname</code> of a custom <code>org.hibernate.connection.ConnectionProvider</code> which provides JDBC connections to Hibernate.</p> <p>e.g. <code>classname.of.ConnectionProvider</code></p>
hibernate.connection.isolation	<p>Sets the JDBC transaction isolation level. Check <code>java.sql.Connection</code> for meaningful values, but note that most databases do not support all isolation levels and some define additional, non-standard isolations.</p> <p>e.g. 1, 2, 4, 8</p>
hibernate.connection.autocommit	Enables autocommit for JDBC pooled connections (it is not recommended). e.g. true false

Property name	Purpose
hibernate.connection.release_mode	<p>Specifies when Hibernate should release JDBC connections. By default, a JDBC connection is held until the session is explicitly closed or disconnected. For an application server JTA datasource, use <code>after_statement</code> to aggressively release connections after every JDBC call. For a non-JTA connection, it often makes sense to release the connection at the end of each transaction, by using <code>after_transaction</code>. <code>auto</code> will choose <code>after_statement</code> for the JTA and CMT transaction strategies and <code>after_transaction</code> for the JDBC transaction strategy.</p> <p>e.g. <code>auto</code> (default) <code>on_close</code> <code>after_transaction</code> <code>after_statement</code></p> <p>This setting only affects Sessions returned from <code>SessionFactory.openSession</code>. For Sessions obtained through <code>SessionFactory.getCurrentSession</code>, the <code>CurrentSessionContext</code> implementation configured for use controls the connection release mode for those Sessions. See Section 2.5, "Contextual sessions"</p>
hibernate.connection.<propertyName>	Pass the JDBC property <code><propertyName></code> to <code>DriverManager.getConnection()</code> .
hibernate.jndi.<propertyName>	Pass the property <code><propertyName></code> to the JNDI <code>InitialContextFactory</code> .

Table 3.5. Hibernate Cache Properties

Property name	Purpose
hibernate.cache.provider_class	<p>The classname of a custom <code>CacheProvider</code>.</p> <p>e.g. <code>classname.of.CacheProvider</code></p>
hibernate.cache.use_minimal_puts	Optimizes second-level cache operation to minimize writes, at the cost of more frequent reads. This setting is most useful for clustered caches and, in Hibernate3, is enabled by default for clustered cache implementations. e.g. <code>true false</code>
hibernate.cache.use_query_cache	Enables the query cache. Individual queries still have to be set cacheable. e.g. <code>true false</code>
hibernate.cache.use_second_level_cache	<p>Can be used to completely disable the second level cache, which is enabled by default for classes which specify a <code><cache></code> mapping.</p> <p>e.g. <code>true false</code></p>
hibernate.cache.query_cache_factory	<p>The classname of a custom <code>QueryCache</code> interface, defaults to the built-in <code>StandardQueryCache</code>.</p> <p>e.g. <code>classname.of.QueryCache</code></p>
hibernate.cache.region_prefix	A prefix to use for second-level cache region names. e.g. <code>prefix</code>
hibernate.cache.use_structured_entries	Forces Hibernate to store data in the second-level cache in a more human-friendly format. e.g. <code>true false</code>

Table 3.6. Hibernate Transaction Properties

Property name	Purpose
---------------	---------

Property name	Purpose
hibernate.transaction.factory_class	<p>The classname of a TransactionFactory to use with Hibernate Transaction API (defaults to JDBCTransactionFactory).</p> <p>e.g. classname.of.TransactionFactory</p>
jta.UserTransaction	<p>A JNDI name used by JTATransactionFactory to obtain the JTA UserTransaction from the application server.</p> <p>e.g. jndi/composite/name</p>
hibernate.transaction.manager_lookup_class	<p>The classname of a TransactionManagerLookup. It is required when JVM-level caching is enabled or when using hilo generator in a JTA environment.</p> <p>e.g. classname.of.TransactionManagerLookup</p>
hibernate.transaction.flush_before_completion	<p>If enabled, the session will be automatically flushed during the before completion phase of the transaction. Built-in and automatic session context management is preferred, see Section 2.5, "Contextual sessions".</p> <p>e.g. true false</p>
hibernate.transaction.auto_close_session	<p>If enabled, the session will be automatically closed during the after completion phase of the transaction. Built-in and automatic session context management is preferred, see Section 2.5, "Contextual sessions".</p> <p>e.g. true false</p>

Table 3.7. Miscellaneous Properties

Property name	Purpose
hibernate.current_session_context_class	<p>Supply a custom strategy for the scoping of the "current" Session. See Section 2.5, "Contextual sessions" for more information about the built-in strategies.</p> <p>e.g. jta thread managed custom.Class</p>
hibernate.query.factory_class	<p>Chooses the HQL parser implementation. e.g. org.hibernate.hql.ast.ASTQueryTranslatorFactory or org.hibernate.hql.classic.ClassicQueryTranslatorFactory</p>
hibernate.query.substitutions	<p>Is used to map from tokens in Hibernate queries to SQL tokens (tokens might be function or literal names, for example). e.g. hqlLiteral=SQL_LITERAL, hqlFunction=SQLFUNC</p>
hibernate.hbm2ddl.auto	<p>Automatically validates or exports schema DDL to the database when the SessionFactory is created. With create-drop, the database schema will be dropped when the SessionFactory is closed explicitly.</p> <p>e.g. validate update create create-drop</p>
hibernate.cglib.use_reflection_optimizer	<p>Enables the use of CGLIB instead of runtime reflection (System-level property). Reflection can sometimes be useful when troubleshooting. Hibernate always requires CGLIB even if you turn off the optimizer. You cannot set this property in hibernate.cfg.xml.</p> <p>e.g. true false</p>

3.4.1. SQL Dialects

Always set the `hibernate.dialect` property to the correct `org.hibernate.dialect.Dialect` subclass for your database. If you specify a dialect, Hibernate will use sensible defaults for some of the other properties listed above. This means that you will not have to specify them manually.

Table 3.8. Hibernate SQL Dialects (`hibernate.dialect`)

RDBMS	Dialect
DB2	<code>org.hibernate.dialect.DB2Dialect</code>
DB2 AS/400	<code>org.hibernate.dialect.DB2400Dialect</code>
DB2 OS390	<code>org.hibernate.dialect.DB2390Dialect</code>
PostgreSQL	<code>org.hibernate.dialect.PostgreSQLDialect</code>
MySQL	<code>org.hibernate.dialect.MySQLDialect</code>
MySQL with InnoDB	<code>org.hibernate.dialect.MySQLInnoDBDialect</code>
MySQL with MyISAM	<code>org.hibernate.dialect.MySQLMyISAMDialect</code>
Oracle (any version)	<code>org.hibernate.dialect.OracleDialect</code>
Oracle 9i	<code>org.hibernate.dialect.Oracle9iDialect</code>
Oracle 10g	<code>org.hibernate.dialect.Oracle10gDialect</code>
Sybase	<code>org.hibernate.dialect.SybaseDialect</code>
Sybase Anywhere	<code>org.hibernate.dialect.SybaseAnywhereDialect</code>
Microsoft SQL Server	<code>org.hibernate.dialect.SQLServerDialect</code>
SAP DB	<code>org.hibernate.dialect.SAPDBDialect</code>
Informix	<code>org.hibernate.dialect.InformixDialect</code>
HypersonicSQL	<code>org.hibernate.dialect.HSQLDialect</code>
Ingres	<code>org.hibernate.dialect.IngresDialect</code>
Progress	<code>org.hibernate.dialect.ProgressDialect</code>
Mckoi SQL	<code>org.hibernate.dialect.MckoiDialect</code>
Interbase	<code>org.hibernate.dialect.InterbaseDialect</code>
Pointbase	<code>org.hibernate.dialect.PointbaseDialect</code>
FrontBase	<code>org.hibernate.dialect.FrontbaseDialect</code>
Firebird	<code>org.hibernate.dialect.FirebirdDialect</code>

3.4.2. Outer Join Fetching

If your database supports ANSI, Oracle or Sybase style outer joins, *outer join fetching* will often increase performance by limiting the number of round trips to and from the database. This is, however, at the cost of possibly more work performed by the database itself. Outer join fetching allows a whole graph of objects connected by many-to-one, one-to-many, many-to-many and one-to-one associations to be retrieved in a single SQL SELECT.

Outer join fetching can be disabled *globally* by setting the property `hibernate.max_fetch_depth` to 0. A setting of 1 or higher enables outer join fetching for one-to-one and many-to-one associations that have been mapped with `fetch="join"`.

See [Section 19.1, "Fetching strategies"](#) for more information.

3.4.3. Binary Streams

Oracle limits the size of byte arrays that can be passed to and/or from its JDBC driver. If you wish to use large instances of binary or serializable type, you should enable `hibernate.jdbc.use_streams_for_binary`. *This is a system-level setting only.*

3.4.4. Second-level and query cache

The properties prefixed by `hibernate.cache` allow you to use a process or cluster scoped second-level cache system with Hibernate. See the [Section 19.2, "The Second Level Cache"](#) for more information.

3.4.5. Query Language Substitution

You can define new Hibernate query tokens using `hibernate.query.substitutions`. For example:

```
hibernate.query.substitutions true=1, false=0
```

This would cause the tokens `true` and `false` to be translated to integer literals in the generated SQL.

```
hibernate.query.substitutions toLowercase=LOWER
```

This would allow you to rename the SQL `LOWER` function.

3.4.6. Hibernate statistics

If you enable `hibernate.generate_statistics`, Hibernate exposes a number of metrics that are useful when tuning a running system via `SessionFactory.getStatistics()`. Hibernate can even be configured to expose these statistics via JMX. Read the Javadoc of the interfaces in `org.hibernate.stats` for more information.

3.5. Logging

Hibernate utilizes [Simple Logging Facade for Java](#) (SLF4J) in order to log various system events. SLF4J can direct your logging output to several logging frameworks (NOP, Simple, log4j version 1.2, JDK 1.4 logging, JCL or logback) depending on your chosen binding. In order to setup logging you will need **slf4j-api.jar** in your classpath together with the jar file for your preferred binding - **slf4j-log4j12.jar** in the case of Log4J. See the SLF4J [documentation](#) for more detail. To use Log4j you will also need to place a **log4j.properties** file in your classpath. An example properties file is distributed with Hibernate in the `src/` directory.

It is recommended that you familiarize yourself with Hibernate's log messages. A lot of work has been put into making the Hibernate log as detailed as possible, without making it unreadable. It is an essential troubleshooting device. The most interesting log categories are the following:

Table 3.9. Hibernate Log Categories

Category	Function
<code>org.hibernate.SQL</code>	Log all SQL DML statements as they are executed
<code>org.hibernate.type</code>	Log all JDBC parameters
<code>org.hibernate.tool.hbm2ddl</code>	Log all SQL DDL statements as they are executed
<code>org.hibernate.pretty</code>	Log the state of all entities (max 20 entities) associated with the session at flush time
<code>org.hibernate.cache</code>	Log all second-level cache activity
<code>org.hibernate.transaction</code>	Log transaction related activity
<code>org.hibernate.jdbc</code>	Log all JDBC resource acquisition
<code>org.hibernate.hql.ast.AST</code>	Log HQL and SQL ASTs during query parsing
<code>org.hibernate.secure</code>	Log all JAAS authorization requests
<code>org.hibernate</code>	Log everything. This is a lot of information but it is useful for troubleshooting

When developing applications with Hibernate, you should almost always work with debug enabled for the category `org.hibernate.SQL`, or, alternatively, the property `hibernate.show_sql` enabled.

3.6. Implementing a NamingStrategy

The interface `org.hibernate.cfg.NamingStrategy` allows you to specify a "naming standard" for database objects and schema elements.

You can provide rules for automatically generating database identifiers from Java identifiers or for processing "logical" column and table names given in the mapping file into "physical" table and column names. This

feature helps reduce the verbosity of the mapping document, eliminating repetitive noise (TBL_ prefixes, for example). The default strategy used by Hibernate is quite minimal.

You can specify a different strategy by calling `Configuration.setNamingStrategy()` before adding mappings:

```
SessionFactory sf = new Configuration()
    .setNamingStrategy(ImprovedNamingStrategy.INSTANCE)
    .addFile("Item.hbm.xml")
    .addFile("Bid.hbm.xml")
    .buildSessionFactory();
```

`org.hibernate.cfg.ImprovedNamingStrategy` is a built-in strategy that might be a useful starting point for some applications.

3.7. XML configuration file

An alternative approach to configuration is to specify a full configuration in a file named `hibernate.cfg.xml`. This file can be used as a replacement for the `hibernate.properties` file or, if both are present, to override properties.

The XML configuration file is by default expected to be in the root of your CLASSPATH. Here is an example:

```
<?xml version='1.0' encoding='utf-8'?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">

<hibernate-configuration>

    <!-- a SessionFactory instance listed as /jndi/name -->
    <session-factory
        name="java:hibernate/SessionFactory">

        <!-- properties -->
        <property name="connection.datasource">java:/comp/env/jdbc/MyDB</property>
        <property name="dialect">org.hibernate.dialect.MySQLDialect</property>
        <property name="show_sql">false</property>
        <property name="transaction.factory_class">
            org.hibernate.transaction.JTATransactionFactory
        </property>
        <property name="jta.UserTransaction">java:comp/UserTransaction</property>

        <!-- mapping files -->
        <mapping resource="org/hibernate/auction/Item.hbm.xml"/>
        <mapping resource="org/hibernate/auction/Bid.hbm.xml"/>

        <!-- cache settings -->
        <class-cache class="org.hibernate.auction.Item" usage="read-write"/>
        <class-cache class="org.hibernate.auction.Bid" usage="read-only"/>
        <collection-cache collection="org.hibernate.auction.Item.bids" usage="read-write"/>

    </session-factory>

</hibernate-configuration>
```

The advantage of this approach is the externalization of the mapping file names to configuration. The `hibernate.cfg.xml` is also more convenient once you have to tune the Hibernate cache. It is your choice to use

either `hibernate.properties` or `hibernate.cfg.xml`. Both are equivalent, except for the above mentioned benefits of using the XML syntax.

With the XML configuration, starting Hibernate is then as simple as:

```
SessionFactory sf = new Configuration().configure().buildSessionFactory();
```

You can select a different XML configuration file using:

```
SessionFactory sf = new Configuration()  
    .configure("catdb.cfg.xml")  
    .buildSessionFactory();
```

3.8. J2EE Application Server integration

Hibernate has the following integration points for J2EE infrastructure:

- » *Container-managed datasources*: Hibernate can use JDBC connections managed by the container and provided through JNDI. Usually, a JTA compatible `TransactionManager` and a `ResourceManager` take care of transaction management (CMT), especially distributed transaction handling across several datasources. You can also demarcate transaction boundaries programmatically (BMT), or you might want to use the optional Hibernate Transaction API for this to keep your code portable.
- » *Automatic JNDI binding*: Hibernate can bind its `SessionFactory` to JNDI after startup.
- » *JTA Session binding*: the Hibernate `Session` can be automatically bound to the scope of JTA transactions. Simply lookup the `SessionFactory` from JNDI and get the current `Session`. Let Hibernate manage flushing and closing the `Session` when your JTA transaction completes. Transaction demarcation is either declarative (CMT) or programmatic (BMT/`UserTransaction`).
- » *JMX deployment*: if you have a JMX capable application server (e.g. JBoss AS), you can choose to deploy Hibernate as a managed MBean. This saves you the one line startup code to build your `SessionFactory` from a `Configuration`. The container will startup your `HibernateService` and also take care of service dependencies (datasource has to be available before Hibernate starts, etc).

Depending on your environment, you might have to set the configuration option `hibernate.connection.aggressive_release` to `true` if your application server shows "connection containment" exceptions.

3.8.1. Transaction strategy configuration

The Hibernate `Session` API is independent of any transaction demarcation system in your architecture. If you let Hibernate use JDBC directly through a connection pool, you can begin and end your transactions by calling the JDBC API. If you run in a J2EE application server, you might want to use bean-managed transactions and call the JTA API and `UserTransaction` when needed.

To keep your code portable between these two (and other) environments we recommend the optional Hibernate Transaction API, which wraps and hides the underlying system. You have to specify a factory class for `Transaction` instances by setting the Hibernate configuration property `hibernate.transaction.factory_class`.

There are three standard, or built-in, choices:

`org.hibernate.transaction.JDBCTransactionFactory`
delegates to database (JDBC) transactions (default)

`org.hibernate.transaction.JTATransactionFactory`
delegates to container-managed transactions if an existing transaction is underway in this context (for example, EJB session bean method). Otherwise, a new transaction is started and bean-managed transactions are used.

`org.hibernate.transaction.CMTTransactionFactory`
delegates to container-managed JTA transactions

You can also define your own transaction strategies (for a CORBA transaction service, for example).

Some features in Hibernate (i.e., the second level cache, Contextual Sessions with JTA, etc.) require access to the JTA `TransactionManager` in a managed environment. In an application server, since J2EE does not standardize a single mechanism, you have to specify how Hibernate should obtain a reference to the `TransactionManager`:

Table 3.10. JTA TransactionManagers

Transaction Factory	Application Server
<code>org.hibernate.transaction.JBossTransactionManagerLookup</code>	JBoss
<code>org.hibernate.transaction.WeblogicTransactionManagerLookup</code>	Weblogic
<code>org.hibernate.transaction.WebSphereTransactionManagerLookup</code>	WebSphere
<code>org.hibernate.transaction.WebSphereExtendedJTATransactionLookup</code>	WebSphere 6
<code>org.hibernate.transaction.OrionTransactionManagerLookup</code>	Orion
<code>org.hibernate.transaction.ResinTransactionManagerLookup</code>	Resin
<code>org.hibernate.transaction.JOTMTransactionManagerLookup</code>	JOTM
<code>org.hibernate.transaction.JOnASTransactionManagerLookup</code>	JOnAS
<code>org.hibernate.transaction.JRun4TransactionManagerLookup</code>	JRun4
<code>org.hibernate.transaction.BESTransactionManagerLookup</code>	Borland ES

3.8.2. JNDI-bound SessionFactory

A JNDI-bound Hibernate `SessionFactory` can simplify the lookup function of the factory and create new Sessions. This is not, however, related to a JNDI bound `Datasource`; both simply use the same registry.

If you wish to have the `SessionFactory` bound to a JNDI namespace, specify a name (e.g. `java:hibernate/SessionFactory`) using the property `hibernate.session_factory_name`. If this property is omitted, the `SessionFactory` will not be bound to JNDI. This is especially useful in environments with a read-only JNDI default implementation (in Tomcat, for example).

When binding the `SessionFactory` to JNDI, Hibernate will use the values of `hibernate.jndi.url`, `hibernate.jndi.class` to instantiate an initial context. If they are not specified, the default `InitialContext` will be used.

Hibernate will automatically place the `SessionFactory` in JNDI after you call `cfg.buildSessionFactory()`. This means you will have this call in some startup code, or utility class in your application, unless you use JMX deployment with the `HibernateService` (this is discussed later in greater detail).

If you use a JNDI `SessionFactory`, an EJB or any other class, you can obtain the `SessionFactory` using a JNDI lookup.

It is recommended that you bind the `SessionFactory` to JNDI in a managed environment and use a static singleton otherwise. To shield your application code from these details, we also recommend to hide the actual lookup code for a `SessionFactory` in a helper class, such as `HibernateUtil.getSessionFactory()`. Note that such a class is also a convenient way to startup Hibernate see chapter 1.

3.8.3. Current Session context management with JTA

The easiest way to handle Sessions and transactions is Hibernate's automatic "current" Session management. For a discussion of contextual sessions see [Section 2.5, "Contextual sessions"](#). Using the "jta" session context, if there is no Hibernate Session associated with the current JTA transaction, one will be started and associated with that JTA transaction the first time you call `sessionFactory.getCurrentSession()`. The Sessions retrieved via `getCurrentSession()` in the "jta" context are set to automatically flush before the transaction completes, close after the transaction completes, and aggressively release JDBC connections after each statement. This allows the Sessions to be managed by the life cycle of the JTA transaction to which it is associated, keeping user code clean of such management concerns. Your code can either use JTA

programmatically through `UserTransaction`, or (recommended for portable code) use the Hibernate Transaction API to set transaction boundaries. If you run in an EJB container, declarative transaction demarcation with CMT is preferred.

3.8.4. JMX deployment

The line `cfg.buildSessionFactory()` still has to be executed somewhere to get a `SessionFactory` into JNDI. You can do this either in a static initializer block, like the one in `HibernateUtil`, or you can deploy Hibernate as a *managed service*.

Hibernate is distributed with `org.hibernate.jmx.HibernateService` for deployment on an application server with JMX capabilities, such as JBoss AS. The actual deployment and configuration is vendor-specific. Here is an example `jboss-service.xml` for JBoss 4.0.x:

```
<?xml version="1.0"?>
<server>

<mbean code="org.hibernate.jmx.HibernateService"
  name="jboss.jca:service=HibernateFactory,name=HibernateFactory">

  <!-- Required services -->
  <depends>jboss.jca:service=RARDeployer</depends>
  <depends>jboss.jca:service=LocalTxCM,name=HsqlDS</depends>

  <!-- Bind the Hibernate service to JNDI -->
  <attribute name="JndiName">java:/hibernate/SessionFactory</attribute>

  <!-- Datasource settings -->
  <attribute name="Datasource">java:HsqlDS</attribute>
  <attribute name="Dialect">org.hibernate.dialect.HSQLDialect</attribute>

  <!-- Transaction integration -->
  <attribute name="TransactionStrategy">
    org.hibernate.transaction.JTATransactionFactory</attribute>
  <attribute name="TransactionManagerLookupStrategy">
    org.hibernate.transaction.JBossTransactionManagerLookup</attribute>
  <attribute name="FlushBeforeCompletionEnabled">true</attribute>
  <attribute name="AutoCloseSessionEnabled">true</attribute>

  <!-- Fetching options -->
  <attribute name="MaximumFetchDepth">5</attribute>

  <!-- Second-level caching -->
  <attribute name="SecondLevelCacheEnabled">true</attribute>
  <attribute name="CacheProviderClass">org.hibernate.cache.EhCacheProvider</attribute>
  <attribute name="QueryCacheEnabled">true</attribute>

  <!-- Logging -->
  <attribute name="ShowSqlEnabled">true</attribute>

  <!-- Mapping files -->
  <attribute name="MapResources">auction/Item.hbm.xml,auction/Category.hbm.xml</attribute>

</mbean>

</server>
```

This file is deployed in a directory called META-INF and packaged in a JAR file with the extension .sar (service archive). You also need to package Hibernate, its required third-party libraries, your compiled persistent classes, as well as your mapping files in the same archive. Your enterprise beans (usually session beans) can be kept in their own JAR file, but you can include this EJB JAR file in the main service archive to get a single (hot-)deployable unit. Consult the JBoss AS documentation for more information about JMX service and EJB deployment.

[Copyright © 2004 Red Hat Middleware, LLC.](#)

[Prev](#)

[Top of page](#)

[Front page](#)

[Next](#)

[Chapter 2. Architecture](#)

[Chapter 4. Persistent Classes](#)